

Softmax

[SOURCE:pdf-softmax]

Softmax converts logits into a probability distribution.

The output probabilities sum to one.

Temperature scaling divides logits by a positive temperature.

A larger temperature produces a flatter distribution.

Gradient descent

[SOURCE:pdf-gradient]

Gradient descent updates parameters opposite to the gradient.

The learning rate controls the size of each update step.

A very large learning rate can cause divergence.

Retrieval evaluation

[SOURCE:pdf-retrieval]

Hit@3 checks whether at least one expected source is in the top three.

MRR is the mean reciprocal rank of the first relevant source.

Source recall measures the fraction of expected sources retrieved.